

# Importing necessary libraries

```
In [2]: import os
import shutil
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import warnings

import tensorflow as tf
from tensorflow.keras import Sequential
from tensorflow.keras.layers import Dense, Dropout, GlobalAveragePooling2D, Conv2D, MaxPool2D, BatchNormalization
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.callbacks import EarlyStopping, LearningRateScheduler
from tensorflow.keras.applications import ResNet50, InceptionV3

from sklearn.model_selection import train_test_split
from sklearn.metrics import precision_score, recall_score, f1_score, accuracy_score
from tabulate import tabulate
warnings.filterwarnings('ignore')
```

## Preparing train and test Image Generator

```
In [3]: dataset_dir = './dataset/cell_images'
train_dir = './dataset/split/train'
val_dir = './dataset/split/validation'
test_dir = './dataset/split/test'

all_data = []
for class_label in ['Parasitized', 'Uninfected']:
    class_path = os.path.join(dataset_dir, class_label)
    for img in os.listdir(class_path):
        all_data.append((os.path.join(class_path, img), class_label))

data_df = pd.DataFrame(all_data, columns=['path', 'label'])
#data_sample = data_df.sample(n=2700, random_state=42)
data_sample = data_df

train_val_data, test_data = train_test_split(data_sample, test_size=0.2, stratify=data_sample['label'], random_state=42)
```

```

train_data, val_data = train_test_split(train_val_data, test_size=0.25, stratify=train_val_data['label'], random_state=42)

def copy_data(data_subset, target_dir):
    for _, row in data_subset.iterrows():
        class_dir = os.path.join(target_dir, row['label'])
        os.makedirs(class_dir, exist_ok=True)
        shutil.copy(row['path'], class_dir)

copy_data(train_data, train_dir)
copy_data(val_data, val_dir)
copy_data(test_data, test_dir)

datagen = ImageDataGenerator(rescale=1/255.0)

width=128
height=128

trainDatagen = datagen.flow_from_directory(
    train_dir,
    target_size=(width, height),
    class_mode='binary',
    batch_size=32
)

valDatagen = datagen.flow_from_directory(
    val_dir,
    target_size=(width, height),
    class_mode='binary',
    batch_size=32
)

testDatagen = datagen.flow_from_directory(
    test_dir,
    target_size=(width, height),
    class_mode='binary',
    batch_size=32,
    shuffle=False
)

```


Found 16758 images belonging to 2 classes.  
 Found 5670 images belonging to 2 classes.  
 Found 5676 images belonging to 2 classes.

In [13]: *# Count the number of images in each category*  
 categories = ['Parasitized', 'Uninfected']

```
counts = [len(os.listdir(os.path.join(dataset_dir, category))) for category in categories]

# Generate the pie chart
plt.figure(figsize=(8, 8))
plt.pie(counts, labels=categories, autopct='%1.1f%%', startangle=90, colors=['#ff9999', '#66b3ff'])
plt.title('Distribution of Uninfected vs Parasitized Cells')
plt.show()
```

## Distribution of Uninfected vs Parasitized Cells



# CNN Model

```
In [4]: cnn_model = Sequential()

cnn_model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(128, 128, 3)))
cnn_model.add(BatchNormalization())
cnn_model.add(MaxPool2D(2, 2))
cnn_model.add(Dropout(0.2))

cnn_model.add(Conv2D(64, (3, 3), activation='relu'))
cnn_model.add(BatchNormalization())
cnn_model.add(MaxPool2D(2, 2))
cnn_model.add(Dropout(0.3))

cnn_model.add(Conv2D(128, (3, 3), activation='relu'))
cnn_model.add(BatchNormalization())
cnn_model.add(MaxPool2D(2, 2))
cnn_model.add(Dropout(0.4))

cnn_model.add(GlobalAveragePooling2D())
cnn_model.add(Dense(128, activation='relu'))
cnn_model.add(Dropout(0.5))
cnn_model.add(Dense(1, activation='sigmoid'))

cnn_model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

early_stop = EarlyStopping(monitor='val_loss', patience=5)

def scheduler(epoch, lr):
    return lr * 0.9 if epoch > 3 else lr

lr_scheduler = LearningRateScheduler(scheduler)

# Train the model
cnn_history = cnn_model.fit(
    x=trainDatagen,
    steps_per_epoch=len(trainDatagen),
    epochs=10,
    validation_data=valDatagen,
    validation_steps=len(valDatagen),
    callbacks=[early_stop, lr_scheduler]
)


# Plot Learning curves
def plotLearningCurve(history):
    epochs = range(1, len(history.history['accuracy']) + 1)
```

```
# Plot accuracy
plt.figure(figsize=(10, 5))
plt.plot(epochs, history.history['accuracy'], label='Training Accuracy')
plt.plot(epochs, history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend(loc='best')
plt.show()


# Plot Loss
plt.figure(figsize=(10, 5))
plt.plot(epochs, history.history['loss'], label='Training Loss')
plt.plot(epochs, history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(loc='best')
plt.show()

plotLearningCurve(cnn_history)
```


Epoch 1/10

524/524  354s 670ms/step - accuracy: 0.7301 - loss: 0.5358 - val\_accuracy: 0.5122 - val\_loss: 2.1460  
- learning\_rate: 0.0010


Epoch 2/10

524/524  167s 319ms/step - accuracy: 0.9398 - loss: 0.1863 - val\_accuracy: 0.7979 - val\_loss: 0.5924  
- learning\_rate: 0.0010


Epoch 3/10

524/524  172s 328ms/step - accuracy: 0.9464 - loss: 0.1666 - val\_accuracy: 0.9229 - val\_loss: 0.2606  
- learning\_rate: 0.0010


Epoch 4/10

524/524  201s 383ms/step - accuracy: 0.9478 - loss: 0.1625 - val\_accuracy: 0.9501 - val\_loss: 0.1520  
- learning\_rate: 0.0010


Epoch 5/10

524/524  174s 332ms/step - accuracy: 0.9516 - loss: 0.1532 - val\_accuracy: 0.9483 - val\_loss: 0.1627  
- learning\_rate: 9.0000e-04


Epoch 6/10

524/524  161s 307ms/step - accuracy: 0.9518 - loss: 0.1486 - val\_accuracy: 0.9563 - val\_loss: 0.1298  
- learning\_rate: 8.1000e-04


Epoch 7/10

524/524  161s 307ms/step - accuracy: 0.9554 - loss: 0.1362 - val\_accuracy: 0.9326 - val\_loss: 0.1947  
- learning\_rate: 7.2900e-04


Epoch 8/10


524/524  161s 307ms/step - accuracy: 0.9573 - loss: 0.1285 - val\_accuracy: 0.9347 - val\_loss: 0.1852  
- learning\_rate: 6.5610e-04

Epoch 9/10


524/524  160s 305ms/step - accuracy: 0.9572 - loss: 0.1267 - val\_accuracy: 0.9330 - val\_loss: 0.1947  
- learning\_rate: 5.9049e-04

Epoch 10/10

524/524  165s 314ms/step - accuracy: 0.9559 - loss: 0.1280 - val\_accuracy: 0.9448 - val\_loss: 0.1467  
- learning\_rate: 5.3144e-04







## CNN Testing

```
In [5]: val_predictions = (cnn_model.predict(testDatagen) > 0.5).astype("int32")
val_true_labels = testDatagen.classes

cnn_accuracy = accuracy_score(val_true_labels, val_predictions)
cnn_precision = precision_score(val_true_labels, val_predictions)
cnn_recall = recall_score(val_true_labels, val_predictions)
cnn_f1 = f1_score(val_true_labels, val_predictions)

print(f"Accuracy: {cnn_accuracy}")
print(f"Precision: {cnn_precision}")
```

```
print(f"Recall: {cnn_recall}")
print(f"F1 Score: {cnn_f1}")
```

178/178 ————— 67s 374ms/step

Accuracy: 0.94538407329105

Precision: 0.9182389937106918

Recall: 0.9777934437786394

F1 Score: 0.9470809149880506

## Resnet Model

```
In [6]: base_resnet = ResNet50(weights='imagenet', include_top=False, input_shape=(128, 128, 3))

for layer in base_resnet.layers[:140]:
    layer.trainable = False
for layer in base_resnet.layers[140:]:
    layer.trainable = True

resnet_model = Sequential([
    base_resnet,
    GlobalAveragePooling2D(),
    Dense(256, activation='relu'),
    Dropout(0.5),
    Dense(1, activation='sigmoid')
])

resnet_model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-4), loss='binary_crossentropy', metrics=['accuracy'])

early_stop = EarlyStopping(monitor='val_loss', patience=5)


def scheduler(epoch, lr):
    return lr * 0.9 if epoch > 5 else lr


lr_scheduler = LearningRateScheduler(scheduler)


resnet_history = resnet_model.fit(
    x=trainDatagen,
    steps_per_epoch=len(trainDatagen),
    epochs=10,
    validation_data=valDatagen,
    validation_steps=len(valDatagen),
    callbacks=[early_stop, lr_scheduler]
)
```


```
# Plot Learning curves
def plotLearningCurve(history, title_prefix="ResNet50"):
    epochs = range(1, len(history.history['accuracy']) + 1)
    # Plot accuracy
    plt.figure(figsize=(10, 5))
    plt.plot(epochs, history.history['accuracy'], label='Training Accuracy')
    plt.plot(epochs, history.history['val_accuracy'], label='Validation Accuracy')
    plt.title(f'{title_prefix} Model Accuracy')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy')
    plt.legend(loc='best')
    plt.show()
    # Plot Loss
    plt.figure(figsize=(10, 5))
    plt.plot(epochs, history.history['loss'], label='Training Loss')
    plt.plot(epochs, history.history['val_loss'], label='Validation Loss')
    plt.title(f'{title_prefix} Model Loss')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.legend(loc='best')
    plt.show()


plotLearningCurve(resnet_history, "ResNet50")
```


Epoch 1/10  
524/524  462s 857ms/step - accuracy: 0.6761 - loss: 0.5996 - val\_accuracy: 0.5012 - val\_loss: 3.0720  
- learning\_rate: 1.0000e-04


Epoch 2/10  
524/524  418s 798ms/step - accuracy: 0.7755 - loss: 0.4647 - val\_accuracy: 0.6663 - val\_loss: 0.6935  
- learning\_rate: 1.0000e-04


Epoch 3/10  
524/524  406s 774ms/step - accuracy: 0.7915 - loss: 0.4363 - val\_accuracy: 0.5235 - val\_loss: 1.7296  
- learning\_rate: 1.0000e-04


Epoch 4/10  
524/524  398s 759ms/step - accuracy: 0.8068 - loss: 0.4094 - val\_accuracy: 0.6090 - val\_loss: 1.3696  
- learning\_rate: 1.0000e-04


Epoch 5/10  
524/524  396s 755ms/step - accuracy: 0.8112 - loss: 0.4048 - val\_accuracy: 0.6961 - val\_loss: 0.5445  
- learning\_rate: 1.0000e-04


Epoch 6/10  
524/524  398s 759ms/step - accuracy: 0.8277 - loss: 0.3790 - val\_accuracy: 0.8076 - val\_loss: 0.3932  
- learning\_rate: 1.0000e-04

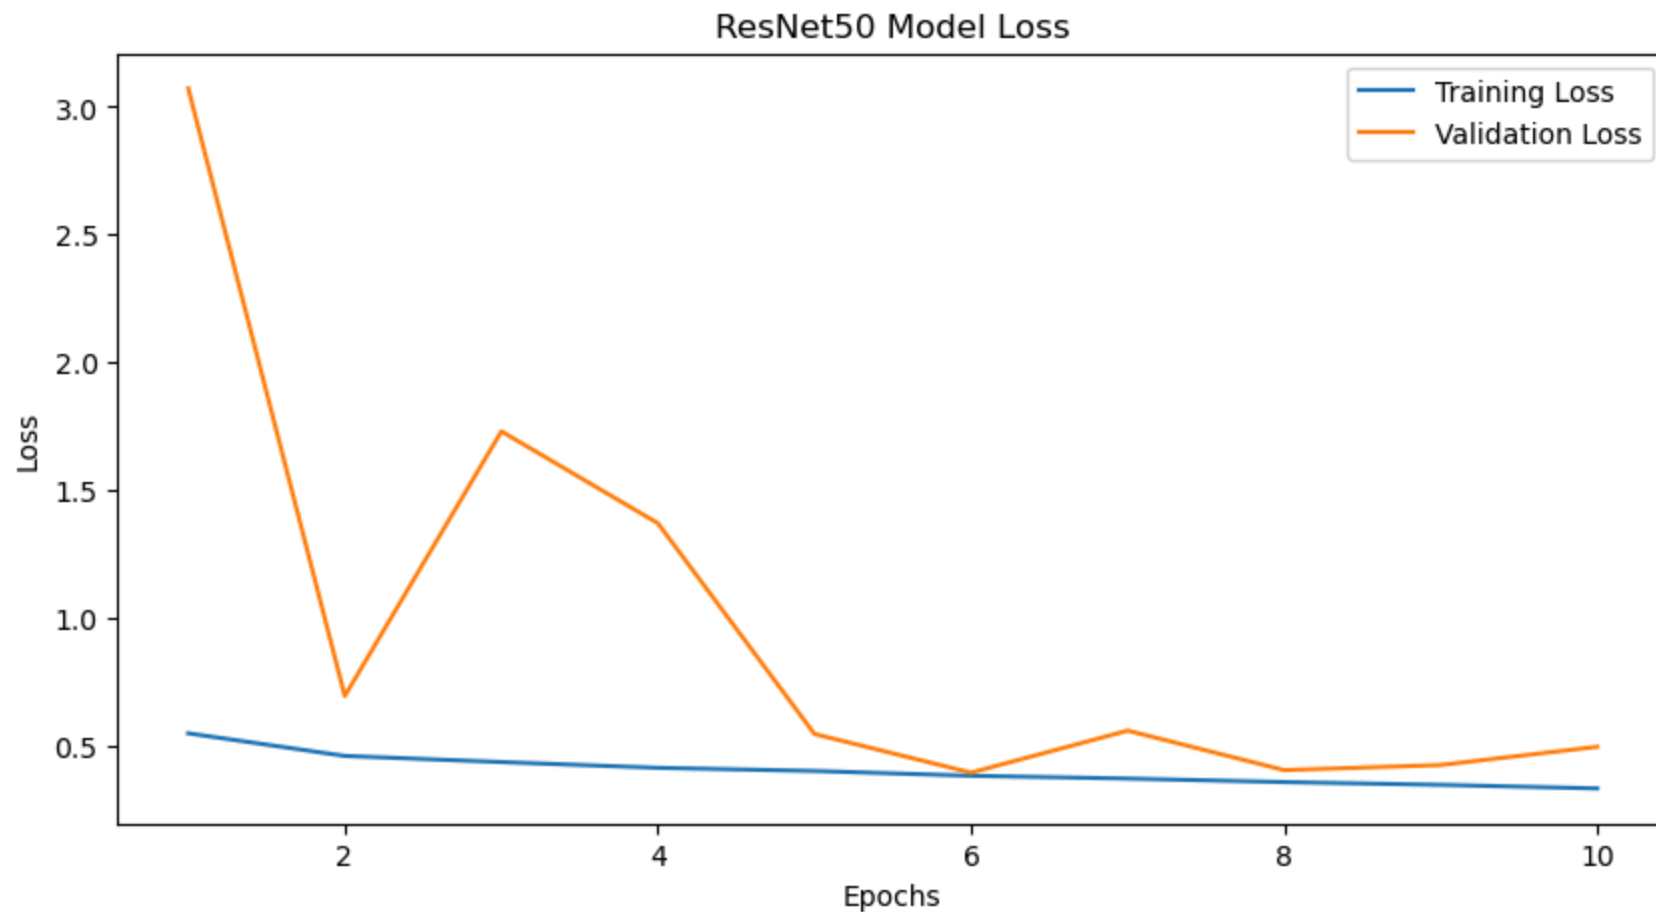
Epoch 7/10  
524/524  408s 778ms/step - accuracy: 0.8296 - loss: 0.3713 - val\_accuracy: 0.6526 - val\_loss: 0.5579  
- learning\_rate: 9.0000e-05

Epoch 8/10  
524/524  400s 763ms/step - accuracy: 0.8404 - loss: 0.3549 - val\_accuracy: 0.8228 - val\_loss: 0.4031  
- learning\_rate: 8.1000e-05

Epoch 9/10  
524/524  401s 765ms/step - accuracy: 0.8446 - loss: 0.3409 - val\_accuracy: 0.7866 - val\_loss: 0.4232  
- learning\_rate: 7.2900e-05

Epoch 10/10  
524/524  399s 763ms/step - accuracy: 0.8499 - loss: 0.3259 - val\_accuracy: 0.7358 - val\_loss: 0.4947  
- learning\_rate: 6.5610e-05





## Resnet Testing

```
In [7]: val_predictions = (resnet_model.predict(testDatagen) > 0.5).astype("int32")
val_true = valDatagen.classes

resnet_accuracy = accuracy_score(val_true_labels, val_predictions)
resnet_precision = precision_score(val_true_labels, val_predictions)
resnet_recall = recall_score(val_true_labels, val_predictions)
resnet_f1 = f1_score(val_true_labels, val_predictions)

print(f"ResNet50 Accuracy: {resnet_accuracy}")
print(f"ResNet50 Precision: {resnet_precision}")
```

```
print(f"ResNet50 Recall: {resnet_recall}")
print(f"ResNet50 F1 Score: {resnet_f1}")
```

178/178 ————— 75s 409ms/step  
 ResNet50 Accuracy: 0.7371388301620859  
 ResNet50 Precision: 0.9153798641136504  
 ResNet50 Recall: 0.522382798731054  
 ResNet50 F1 Score: 0.6651705565529623

## Inception V3 Model

```
In [8]: base_inception = InceptionV3(weights='imagenet', include_top=False, input_shape=(128, 128, 3))

for layer in base_inception.layers[:249]:
    layer.trainable = False
for layer in base_inception.layers[249:]:
    layer.trainable = True

inception_model = Sequential([
    base_inception,
    GlobalAveragePooling2D(),
    Dense(256, activation='relu'),
    Dropout(0.5),
    Dense(1, activation='sigmoid')
])

inception_model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-4), loss='binary_crossentropy', metrics=['accuracy'])

early_stop = EarlyStopping(monitor='val_loss', patience=3)

inception_history = inception_model.fit(
    x=trainDatagen,
    steps_per_epoch=len(trainDatagen),
    epochs=10,
    validation_data=valDatagen,
    validation_steps=len(valDatagen),
    callbacks=[early_stop]
)

# Plot Learning curves
def plotLearningCurve(history, title_prefix="InceptionV3"):
    epochs = range(1, len(history.history['accuracy']) + 1)
    # Plot accuracy
```

```

plt.figure(figsize=(10, 5))
plt.plot(epochs, history.history['accuracy'], label='Training Accuracy')
plt.plot(epochs, history.history['val_accuracy'], label='Validation Accuracy')
plt.title(f'{title_prefix} Model Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend(loc='best')
plt.show()
# Plot Loss
plt.figure(figsize=(10, 5))
plt.plot(epochs, history.history['loss'], label='Training Loss')
plt.plot(epochs, history.history['val_loss'], label='Validation Loss')
plt.title(f'{title_prefix} Model Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(loc='best')
plt.show()

plotLearningCurve(inception_history, "InceptionV3")

```

Epoch 1/10

**524/524** ————— 233s 419ms/step - accuracy: 0.8792 - loss: 0.2962 - val\_accuracy: 0.9395 - val\_loss: 0.1575

Epoch 2/10

**524/524** ————— 223s 426ms/step - accuracy: 0.9532 - loss: 0.1275 - val\_accuracy: 0.9423 - val\_loss: 0.1618


Epoch 3/10


**524/524** ————— 219s 418ms/step - accuracy: 0.9707 - loss: 0.0791 - val\_accuracy: 0.9411 - val\_loss: 0.1912

Epoch 4/10

**524/524** ————— 220s 420ms/step - accuracy: 0.9826 - loss: 0.0507 - val\_accuracy: 0.9379 - val\_loss: 0.2072







## Inception V3 Testing

```
In [9]: val_predictions = (inception_model.predict(testDatagen) > 0.5).astype("int32")
val_true_labels = testDatagen.classes

inception_accuracy = accuracy_score(val_true_labels, val_predictions)
inception_precision = precision_score(val_true_labels, val_predictions)
inception_recall = recall_score(val_true_labels, val_predictions)
inception_f1 = f1_score(val_true_labels, val_predictions)

print(f"InceptionV3 Accuracy: {inception_accuracy}")
print(f"InceptionV3 Precision: {inception_precision}")
```

```
print(f"InceptionV3 Recall: {inception_recall}")
print(f"InceptionV3 F1 Score: {inception_f1}")
```

178/178 ————— 42s 222ms/step  
 InceptionV3 Accuracy: 0.9342847075405215  
 InceptionV3 Precision: 0.9095744680851063  
 InceptionV3 Recall: 0.9643990130419458  
 InceptionV3 F1 Score: 0.9361847733105219

```
In [10]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay


def plot_confusion_matrix(test_generator, model, model_name):
    # Predict the classes
    predictions = (model.predict(test_generator) > 0.5).astype("int32")
    true_labels = test_generator.classes

    # Compute confusion matrix
    cm = confusion_matrix(true_labels, predictions)
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=test_generator.class_indices.keys())


    # Plot the confusion matrix
    plt.figure(figsize=(8, 8))
    disp.plot(cmap=plt.cm.Blues, ax=plt.gca())
    plt.title(f'Confusion Matrix for {model_name}')
    plt.savefig(f'{model_name}_confusion_matrix.png')
    plt.show()

plot_confusion_matrix(testDatagen, cnn_model, "CNN")
plot_confusion_matrix(testDatagen, resnet_model, "ResNet")
plot_confusion_matrix(testDatagen, inception_model, "InceptionV3")
```


178/178 ————— 10s 56ms/step



178/178 ————— 39s 216ms/step



178/178 ————— 20s 113ms/step



```
In [12]: model_scores = {  
    "Model": ["CNN", "ResNet50", "InceptionV3"],  
    "Accuracy": [cnn_accuracy, resnet_accuracy, inception_accuracy],  
    "Precision": [cnn_precision, resnet_precision, inception_precision],  
    "Recall": [cnn_recall, resnet_recall, inception_recall],  
}
```

```
"F1 Score": [cnn_f1, resnet_f1, inception_f1]
}

scores_df = pd.DataFrame(model_scores)

print(tabulate(scores_df, headers="keys", tablefmt="fancy_grid", floatfmt=".2f"))
```

	Model	Accuracy	Precision	Recall	F1 Score
0	CNN	0.95	0.92	0.98	0.95
1	ResNet50	0.74	0.92	0.52	0.67
2	InceptionV3	0.93	0.91	0.96	0.94

Conclusion: CNN and InceptionV3 are the best models while CNN is slightly better than InceptionV3. ResNet50 needs a lot of improvements if it is to be effective on this dataset.

In [ ]: